

¹ Pyfive: A pure-python HDF5 reader

² **Bryan Lawrence**  ¹, **Ezequiel Cimadevilla**  ², **Wout De Nolf**  ⁶, **David Hassell**  ¹, **Jonathan Helmus**³, **Brian Maranville**  ⁴, **Kai Mühlbauer**  ⁵, and ⁴ **Valeriu Predoi** 

⁵ 1 NCAS-CMS, Meteorology Department, University of Reading, Reading, United Kingdom.  ²
⁶ Instituto de Física de Cantabria (IFCA), CSIC-Universidad de Cantabria, Santander, Spain. ³ TBD ⁴
⁷ NIST Center for Neutron Research ⁵ Institute of Geosciences, Meteorology Section, University of Bonn,
⁸ Germany. ¹ European Synchrotron Radiation Facility (ESRF), Grenoble, France.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- ⁹ [Review](#) 
- ¹⁰ [Repository](#) 
- ¹¹ [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- ¹² [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).
²¹
²²
²³

Summary

Pyfive (<https://pyfive.readthedocs.io/en/latest/>) is an open-source thread-safe pure Python package for reading data stored in HDF5. While it is not a complete implementation of all the specifications and capabilities of HDF5, it includes all the core functionality necessary to read gridded datasets, whether stored contiguously or with chunks, and to carry out the necessary decompression for the standard options. All data access is fully lazy, the data is only read from storage when the numpy data arrays are manipulated. Originally developed some years ago, the package has recently been upgraded to support lazy access, and to add missing features necessary for handling all the environmental data known to the authors. It is now a realistic option for production data access in environmental science and more widely. The API is based on that of h5py (which is a python shimmy over the HDF5 c-library, and hence is not thread-safe), with some API extensions to help optimise remote access. With these extensions, coupled with thread safety, many of the limitations precluding the efficient use of HDF5 (and NetCDF4) on cloud storage have been removed.

Statement of need

HDF5 is probably the most important data format in physical science, used across the piste. It is particularly important in environmental science, particularly given the fact that NetCDF4 is HDF5 under the hood. From satellite missions, to climate models and radar systems, the default binary format has been HDF5 for decades. While newer formats are starting to get mindshare, there are petabytes, if not exabytes of existing HDF5, and there are still many good use-cases for creating new data in HDF5. However, despite the history, there are few libraries for reading HDF5 file data that do not depend on the official HDF5 library maintained by the HDFGroup, and in particular, apart from pyfive, in Python there are none that cover the needs of environmental science. While the HDF5 c-library is reliable and performant, and battle-tested over decades, there are some caveats to depending upon it: Firstly, it is not thread-safe, secondly, the code is large and complex, and should anything happen to the financial stability of The HDF5group, it is not obvious the C-code could be maintained. Finally, the code complexity also meant that it was not suitable for developing bespoke code for data recovery in the case of partially corrupt data. From a long-term curation perspective both of these last two constraints are a concern.

The original implementation of pyfive (by JH), which included all the low-level functionality to deal with the internals of an HDF5 file was developed with POSIX access in mind. The recent upgrades were developed with the use-case of performant remote access to curated data as the primary motivation, but with additional motivations of having a lightweight HDF5

43 reader capable of deploying in resource or operating-system constrained environments (such
44 as mobile), and one that could be maintained long-term as a reference reader for curation
45 purposes. The lightweight deployment consequences of a pure-python HDF5 reader need no
46 further introduction, but as additional motivation we now expand on the issues around remote
47 access and curation.

48 Thread safety has become a concern given the wide use of Dask in python based analysis
49 workflows, and this, coupled with a lack of user knowledge about how to efficiently use HDF5,
50 has led to a community perception that HDF5 is not fit for remote access (especially on cloud
51 storage). Issues with thread safety arise from the underlying HDF5 c-library, and cannot be
52 resolved in any solution depending on that library, hence the desire for a pure python solution.
53 Remote access has been bedevilled by the widespread need to access remotely data which has
54 been chunked and compressed, combined with the use of HDF5 data which was left in the
55 state it was when the data was produced - often with default unsuitable chunking and with
56 interleaved chunk indexes and data. Solutions have mainly consisted of reformatting the data
57 (and rechunking it at the same time) or utilising kerchunk mediated direct access to chunked
58 HDF5 data (<https://fsspec.github.io/kerchunk/>). However, in practice using kerchunk requires
59 the data provider to generate kerchunk indices to support remote users, and it leads to issues
60 of synchronicity between indices and changing datasets.

61 This version of pyfive was developed with these use-cases in mind. There is now full support
62 for lazy loading of chunked data, and methods are provided to give users all the benefits of
63 using kerchunk, but without the need for a priori generation. Because pyfive can access and
64 cache (in the client) the b-tree (index) on a variable-by-variable basis, most of the benefits of
65 kerchunk are gained without any of the constraints. However, the kerchunk index is always a
66 contiguous object accessible with one get transaction, this is not necessarily the case with the
67 b-tree, unless the source data has been repacked to ensure contiguous metadata using a tool
68 like h5repack.

69 However, much of the community is unaware of the possibility of repacking the index metadata,
70 and this together with relatively opaque information about the internal structure of files (and
71 hence the necessity or other wise of such repacking), means that repacking is rarely done. To
72 help with this process, pyfive also includes extensions to expose information about how data
73 and indexes are distributed in the files. With these tools, index extraction with pyfive can be
74 comparable in performance to obtaining a kerchunk index, and completely opaque to the user.

75 With the use of pyfive, suitably repacked and rechunked HDF5 data can now be considered
76 "cloud-optimised", insofar as with lazy loading, improved index handling, and thread-safety,
77 there are no "format-induced" constraints on performance during remote access. To aid in
78 discovering whether or not a given HDF5 dataset is cloud-optimised, pyfive also now provides
79 a simple method to find out.

80 The issues of the dependency on a complex code maintained by one private company in the
81 context of maintaining data access (over decades, and potentially centuries), can only be
82 mitigated by ensuring that the data format is well documented, that data writers use only the
83 documented features, and that public code exists which can be relatively easily maintained.
84 The HDF5group have provided good documentation for the core features of HDF5 which
85 include all those of interest to the weather and climate community who motivated this reboot
86 of pyfive, and while there is a community of developers beyond the HDF5 group (including
87 some at the publicly funded Unidata institution), recent events suggest that given most of
88 those developers and their existing funding are US based, some spreading of risk would be
89 desirable. To that end, a pure-python code, which is relatively small and maintained by an
90 international constituency, alongside the existing c-code, provides some assurance that the
91 community can maintain HDF5 access for the foreseeable future. A pure python code also
92 makes it easier to develop scripts which can work around data and metadata damage should
93 they occur.

94 Examples

95 Remote Access

96 A notable feature of the recent pyfive upgrade is that it was carried out with thread-safety and
97 remote access using fsspec (filesystem-spec.readthedocs.io) in mind. We provide two examples
98 of using pyfive to access remote data, one in S3, and one behind a modern http web server:

99 For accessing the data on S3 storage, we will have to set up an s3fs virtual file system, then
100 pass it to Pyfive:

```
import pyfive
import s3fs
```

```
# storage options for an anon S3 bucket
storage_options = {
    'anon': True,
    'client_kwargs': {'endpoint_url': "https://s3server.ac.uk"}
}
fs = s3fs.S3FileSystem(**storage_options)
file_uri = "s3-bucket/myfile.nc"
with fs.open(file_uri, 'rb') as s3_file:
    nc = pyfive.File(s3_file)
    dataset = nc[var]
```

101 for an HTTPS data server, the usage is similar:

```
import fsspec
import pyfive

fs = fsspec.filesystem('http')
with fs.open("https://site.com/myfile.nc", 'rb') as http_file:
    nc = pyfive.File(http_file)
    dataset = nc[var]
```

102 Cloud Optimisation

103 To be fully cloud optimised an HDF5 file needs to have a contiguous index for each variable,
104 and the chunks for each variable need to be broadly contiguous within the file.

105 When these criteria are met, indexes can be read efficiently, and middleware such as fsspec
106 can make sensible use of readahead caching strategies.

107 HDF5 data files direct from simulations and instruments are often not in this state as information
108 about the number of variables, the number of chunks per variable, and the compressed size
109 of those variables is not known as the data is being produced. In such cases the data is also
110 often not chunked along the dimensions being added to as the file is written (since it would
111 have to be buffered first).

112 Of course, once the file is produced, such information is available. Metadata can be repacked to
113 the front of the file and variables can be rechunked and made continuous - which is effectively
114 the same process undertaken when HDF5 data is reformatted to other “so-called” cloud
115 optimised formats.

116 The HDF5 library provides a tool “h5repack” which can do this, provided it is driven with
117 suitable information about required chunk shape and the expected size of metadata fields. This

¹¹⁸ version of pyfive supports both method to query whether such repacking is necessary, and to
¹¹⁹ extract necessary parameters.

EXAMPLES

¹²⁰ Citations

¹²¹ Citations to entries in paper.bib should be in [rMarkdown](#) format.

¹²² If you want to cite a software repository URL (e.g. something on GitHub without a preferred
¹²³ citation) then you can do it with the example BibTeX entry below for Smith et al. ([2020](#)).

¹²⁴ For a quick reference, the following citation commands can be used: - @author:2001 ->
¹²⁵ "Author et al. (2001)" - [@author:2001] -> "(Author et al., 2001)" - [@author1:2001;
¹²⁶ @author2:2001] -> "(Author1 et al., 2001; Author2 et al., 2002)"

¹²⁷ Figures

¹²⁸ Figures can be included like this: Caption for example figure. and referenced from text using
¹²⁹ [section](#).

¹³⁰ Figure sizes can be customized by adding an optional second parameter: Caption for example
¹³¹ figure.

¹³² Acknowledgements

¹³³ Most of the developments outlined here that have occurred since V0.5 (primarily authored
¹³⁴ by JH) have been supported by the UK Met Office and UKRI via 1) UK Excalibur Exascale
¹³⁵ programme (project ExcaliWork), 2) the UKRI Digital Research Infrastructure programme
¹³⁶ (project WacaSoft), and 3) the long term single centre science programme of the UK National
¹³⁷ Center for Atmospheric Science (NCAS). Ongoing maintenance of pyfive is expected to continue
¹³⁸ under the auspices of NCAS national capability funding.

¹³⁹ References

¹⁴⁰ Smith, A. M., Thaney, K., & Hahnel, M. (2020). Fidgit: An ungodly union of GitHub and
¹⁴¹ figshare. In *GitHub repository*. GitHub. <https://github.com/arfon/fidgit>